

Strange Evals

Why?

- ▶ Benchmark scores give a ten-thousand foot view of capability; this group is all about building more powerful and nuanced intuitions of capability

“Become one with the data.”


— Andrej Karpathy, “A Recipe for Training Neural Networks,” 2019 [10]

The State of VLM Benchmarks

“While large language models with vision capabilities (VLMs), e.g., GPT-4o and Gemini-1.5 Pro, score high on many vision-understanding benchmarks, they are still struggling with low-level vision tasks that are easy to humans.”

— Rahmanzadehgervi et al., “Vision Language Models Are Blind,” 2024 [9]

The VLMS



Benchmark Usage

	MMMU Pro	MathVista	A2D	CharXv	MMMU	Video-MME	RefCOCO	ChartQA	DocVQA	LVBench	OCRBench	DynaMath	MLVU	MMStar	MTBench	RealWorldQA	VideoMMMU	BUNK	Hallusion	MUJBENCH	ZenBench	MMBench-de	SimpleVQA
GPT-5.5	81.2	76.0	96.2	—	—	71.2	—	92.1	91.5	—	—	—	—	—	—	—	—	—	—	—	—	—	—
Claude Opus 4.7	81.4	79.3	94.8	82.1	84.1	67.8	—	88.0	93.0	—	—	—	—	—	—	—	—	—	—	—	—	—	—
Qwen3.6-35B-A3B	75.3	86.4	92.7	78.0	81.7	82.5	92.0	—	—	71.4	89.4	—	86.2	—	74.6	85.3	83.7	—	69.8	—	34.4	92.8	58.9
Gemma-4-31B	76.9	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
GLM-4.6V	66.0	85.2	88.8	63.2	76.0	—	88.6	—	—	59.5	86.5	54.5	—	75.9	—	—	74.7	65.5	—	77.1	25.8	—	—
InternVL3.5	—	82.7	87.3	48.3	77.7	72.9	92.4	88.0	94.9	67.1	907	46.5	78.2	77.9	76.5	75.2	—	61.4	57.3	47.5	—	—	—


[4, 5, 6, 7, 8, 11]

Example

The stated goal of MMMU, the most cited vision benchmark, is to understand how close a model is to performing at the same level as “skilled adults”.

[1]

MMMU Pro – 5 Random Examples



The diagram above depicts iron filings sprinkled around three permanent magnets. Pole R is the same pole as

- A) Y and Z
- B) X and T
- C) R and T
- D) X and Y
- E) T and Z
- F) T and Y
- G) R and Z
- H) R and Y
- I) X and Y



What is your opinion on the fact that certain large insects act as predators of pests? Are they allies or enemies in your view?

- A) Neutral; it doesn't affect the crops positively or negatively.
- B) Symbiotic; it has a mutual relationship with the crops.
- C) Vector; it is not a pest directly but spreads viruses which damage yield.
- D) Threat; it may not harm the crops directly, but its presence could attract other pests.
- E) Beneficial; it helps in biological control of pests.
- F) Neither; this large insect is not a problem, nor is it a beneficial.
- G) I don't know and don't want to guess.
- H) Friend; this is clearly a predator of pests.
- I) Foe; this is a pest of crops.
- J) Parasitic; it feeds on the pests but also harms the crops.



Refer to the figure, which term best describes the practice where members create the visual essence of a chapter on a three-foot strip of paper, discuss it, and then assemble the strips into the shape of sun rays?

- A) Chapter Ray
- B) Story Ray
- C) Narrative Ray
- D) Magical Realism Ray
- E) Plot Ray
- F) Animal Fantasy Ray
- G) Folklore Ray
- H) Trickster Tale Ray
- I) Fiction Ray
- J) Fable Ray

Test results	Number of absences	Number of deaths from COVID
never out of bed	20,076	756
once a month	140,386	2,041
2-3 times a month	143,183	1,835
4-6 times a month	42,319	270
every day	81,386	452
Total	325,950	5,354

The overall Chi-square value was


- A) 1196.29
- B) 1700.21
- C) 1300.54
- D) 1567.34
- E) 1082.76
- F) 1552.34
- G) 1203.82
- H) 1100.65
- I) 987.45
- J) 1356.98

Project	Beta	IRR
W	.80	9.4%
X	.95	10.9
Y	1.15	13.0
Z	1.45	14.2

An all-equity firm is considering the following projects: The T-bill rate is 3.5 percent, and the expected return on the market is 11 percent. Which projects would be incorrectly accepted or rejected if the firm's overall cost of capital was used as a hurdle rate?

- A) Project Z would be incorrectly rejected; Project Y would be incorrectly accepted.
- B) Project Z would be incorrectly accepted; Project X would be incorrectly rejected.
- C) Project X would be incorrectly rejected; Project Y would be incorrectly accepted.
- D) Project X would be incorrectly accepted; Project Y would be incorrectly rejected.
- E) No project would be incorrectly rejected; Project X would be incorrectly accepted.
- F) Project Y would be incorrectly rejected; Project X would be incorrectly accepted.
- G) No project would be incorrectly rejected; Project Z would be incorrectly accepted.
- H) Project X would be incorrectly rejected; Project Z would be incorrectly accepted.
- I) Project Y would be incorrectly rejected; Project Z would be incorrectly accepted.
- J) Project Z would be incorrectly rejected; Project X would be incorrectly accepted.

MMMU Pro – Answers



The diagram above depicts iron filings sprinkled around three permanent magnets. Pole R is the same pole as

- A) Y and Z
- B) X and T
- C) R and T
- D) X and Y
- E) T and Z
- F) T and Y
- G) R and Z
- H) R and Y
- I) X and Z



What is your opinion on the fact that certain large insects act as predators of pests? Are they allies or enemies in your view?

- A) Neutral; it doesn't affect the crops positively or negatively.
- B) Symbiotic; it has a mutual relationship with the crops.
- C) Vector; it is not a pest directly but spreads viruses which damage yield.
- D) Threat; it may not harm the crops directly, but its presence could attract other pests.
- E) Beneficial; it helps in biological control of pests.
- F) Neither; this large insect is not a problem, nor is it a beneficial.
- G) I don't know and don't want to guess.
- H) Friend; this is clearly a predator of pests.
- I) Foe; this is a pest of crops.
- J) Parasitic; it feeds on the pests but also harms the crops.



Refer to the figure, which term best describes the practice where members create the visual essence of a chapter on a three-foot strip of paper, discuss it, and then assemble the strips into the shape of sun rays?

- A) Chapter Ray
- B) Story Ray
- C) Narrative Ray
- D) Magical Realism Ray
- E) Plot Ray
- F) Animal Fantasy Ray
- G) Folklore Ray
- H) Trickster Tale Ray
- I) Fiction Ray
- J) Fable Ray

Test results	Number of absences	Number of deaths from COVID
none over a full	10,000	100
once a month	150,000	1,000
2-3 times a month	140,000	1,000
4-6 times a month	40,000	1,000
every day	10,000	400
Total	400,000	4,000

The overall Chi-square value was


- A) 1196.29
- B) 1700.21
- C) 1300.54
- D) 1567.34
- E) 1082.76
- F) 1552.34
- G) 1203.82
- H) 1100.65
- I) 987.45
- J) 1356.98

Project	Beta	IRR
W	.80	9.4%
X	.95	10.9
Y	1.15	13.0
Z	1.45	14.2

An all-equity firm is considering the following projects: The T-bill rate is 3.5 percent, and the expected return on the market is 11 percent. Which projects would be incorrectly accepted or rejected if the firm's overall cost of capital was used as a hurdle rate?


- A) Project Z would be incorrectly rejected; Project Y would be incorrectly accepted.
- B) Project Z would be incorrectly accepted; Project X would be incorrectly rejected.
- C) Project X would be incorrectly rejected; Project Y would be incorrectly accepted.
- D) Project X would be incorrectly accepted; Project Y would be incorrectly rejected.
- E) No project would be incorrectly rejected; Project X would be incorrectly accepted.
- F) Project Y would be incorrectly rejected; Project X would be incorrectly accepted.
- G) No project would be incorrectly rejected; Project Z would be incorrectly accepted.
- H) Project X would be incorrectly rejected; Project Z would be incorrectly accepted.
- I) Project Y would be incorrectly rejected; Project Z would be incorrectly accepted.
- J) Project Z would be incorrectly rejected; Project X would be incorrectly accepted.

Cherry Picked ZeroBench Example




Question: What percentage of the geese are oriented such that their body faces south on an 8-point compass? Give your answer to one decimal place.

MMMU Error Analysis (GPT-4V)



“We meticulously examine 150 randomly sampled error instances from GPT-4V’s predictions. These instances are analyzed by expert annotators who identify the root causes of mispredictions based on their knowledge and the golden explanations if available.”

MMMU Pro vs ZeroBench (Sub)



VLMs are Blind





Fig. 1: VLMs fail on the simple tasks of BlindTest.

VLMs are Blind




VLMs are Blind


Benchmark performance

To highlight visual reasoning improvement versus our previous multimodal models, we tested OpenAI o3 and o4-mini on a diverse set of human exams and ML benchmarks. These new visual reasoning models significantly outperform their predecessors on **all** multimodal tasks we tested.


problem-solving



MathVista
Visual math reasoning



VLMs are blind
Visual perception primitives



VLMs are Biased

a. 🐕



	2	✗
	2	✗
	4	✗
	2	✗
	2	✗
Bias	2	✗
GT	3	✓

How many legs does this animal have?

References

1. Yue et al., "MMMU: A Massive Multi-discipline Multimodal Understanding and Reasoning Benchmark for Expert AGI," CVPR 2024.
2. Roberts et al., "ZeroBench: An Impossible Visual Benchmark for Contemporary Large Multimodal Models," arXiv:2502.09696, 2025.
3. THUDM, "GLM-4.5V Technical Report," 2025.
4. Qwen Team, "Qwen3.6-35B-A3B: Agentic Coding Power, Now Open to All," 2026.
5. OpenAI, "Introducing GPT-5.5," 2026.
6. Google DeepMind, "Gemma 4 Technical Report," 2026.
7. THUDM, "GLM-4.6V Technical Report," 2025.
8. OpenGVLab, "InternVL3.5 Technical Report," 2025.
9. Rahmanzadehgervi et al., "Vision Language Models Are Blind," arXiv:2407.06581, 2024.
10. Karpathy, "A Recipe for Training Neural Networks," <https://karpathy.github.io/2019/04/25/recipe/>, 2019.
11. Anthropic, "Introducing Claude Opus 4.7," 2026.
12. OpenAI, "Introducing o3 and o4-mini," 2025.
13. Vo et al., "Vision Language Models are Biased," arXiv:2505.23941, 2025.

Appendix: Links

- ▶ Proof the sun-ray task actually exists with that phrasing